

Both files are substantially rewritten. I re-verified the arithmetic and diffed them: P1 and P6 are now byte-identical between variants, and divergence is confined to three adapter implementations. Nearly every blocker from the prior review is closed.

1. Headline grades

|                                                 | Open-source | LocalStack |
|-------------------------------------------------|-------------|------------|
| Overall program grade                           | A–          | B+         |
| Design quality (ignoring external dependencies) | A–          | A–         |
| Launch readiness                                | B–          | C+         |

The LocalStack grade is capped by its own line 38 (unresolved Hobby-terms confirmation), not by design quality. Both share one gating constraint: every rubric, floor, and supplied asset named in "Production-readiness dependencies" is still unbuilt.

2. Summary evaluation matrix

| Evaluated dimension            | OS | LS | Assessment                                                                                                                                                                                                                                                                                               |
|--------------------------------|----|----|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Pacing & workload (440h)       | A– | B+ | Arithmetic exact on all six projects and the total. Hidden hours genuinely removed — preflight at admission, pinned reduced profiles, supplied scaffolds, 5 recordings instead of ~11. Residual risk sits in three tasks, not across the program.                                                        |
| Internal consistency           | A  | A  | No contradictions between summary tables, delivery contract, task bodies, and workload table. The fde-focused-6-project-tree.md file is now stale against both.                                                                                                                                          |
| Technical relevance (2026)     | A  | A  | OTel/Tempo, pgvector+FTS hybrid with RRF and cross-encoder rerank, circuit breaker + DLQ, STRIDE/OWASP-LLM, LangGraph, deterministic evals, IMDSv2/least privilege, ports & adapters. Correctly excludes fine-tuning and K8s — matches the 0-of-41-postings evidence.                                    |
| FDE field competencies         | A  | A  | Four discovery reps, SoW with exclusions, two change-control reps, incident comms, executive decision, adoption, handback, teardown, playbook, context-switch inject, and now undocumented-legacy integration. Only the live drill is missing.                                                           |
| Local-first ergonomics         | A– | B+ | 16 GB floor, 10 GB container budget, ≤3B quantized model, emulator fallback, admission preflight. Genuinely engineered rather than asserted. LS carries one more container plus a third-party account.                                                                                                   |
| AI integration rigor           | A  | A  | AI is treated as an untrusted component: schema validation, step/token/tool budgets, human checkpoint, audit trail, injection controls, output validation. Judge and Ragas are explicitly kept out of CI with variance reported — this is better than most industry practice.                            |
| Assessment determinism         | A– | A– | Fixed transcripts, checkpoint-triggered injects, published floors, chatbot demoted to unassessed. Docked because the floors and rubric anchors do not exist yet, and 6.5's accept/reject decision has no published derivation rule (6.3 does).                                                           |
| Evidence & acceptance design   | B+ | B+ | Floors added in 2.4, 3.3, 5.2 and thresholds in 6.3. P1, P4, 6.4 and 6.5 still have named artifacts but no pass condition. Largest remaining authoring gap.                                                                                                                                              |
| Repository & continuity        | A  | A  | The checkpoint model is the strongest single addition: it kills defect propagation, caps integration hours, protects grading integrity, and keeps five physical repos — preserving the dual-track economics. One consequence is unstated (see §5).                                                       |
| Cloud & budget governance      | A  | A  | Fixed profile, prohibited-resource list, constrained role and templates, budget actions, supplied DNS/TLS/SSM, automated teardown, verifier, \$80 target inside a \$180 cap with one remediation, endpoint by end of Week 20, 3-day SLA, no GPU. Now tighter than the canonical runbook on window logic. |
| Claim honesty / governance     | A  | A  | Explicit evidence boundary, per-row "what this does not prove", fidelity notes in 2.2/3.2/4.3, pressure inject in 4.5. Exemplary discipline.                                                                                                                                                             |
| Tooling economy                | A– | B+ | Removed OpenSearch, OpenBao, DynamoDB, KMS, IAM, SNS, Lambda/EventBridge/API Gateway, GPU, and every either/or fork. LS keeps a container whose payoff shrank: SDK code now lives inside supplied adapters, so students barely write AWS calls.                                                          |
| Cross-variant parity           | A  | A  | ~27 of 31 tasks identical; P1 and P6 byte-identical; skills lines now differ correctly where tools differ.                                                                                                                                                                                               |
| Admission gate                 | B  | B  | Added async behavior, unfamiliar-codebase debugging, written communication, preflight. Still no production-history bar despite the market's four-year median and "shipped at scale, verified at intake".                                                                                                 |
| Portfolio & career value       | A– | A– | Five named artifacts, public index, sanitized repo, reference architecture, playbook. Docked for the reference-code point and the ungraded demo day.                                                                                                                                                     |
| Marketing-claim safety         | A  | A  | "Multi-Agent" gone, fine-tuning out of scope, status header present, no placement or salary language. Safe to hand to marketing with the evidence boundary attached.                                                                                                                                     |
| Production-authoring readiness | B– | C+ | Dependencies honestly listed at line 23 — and all of them are unbuilt: scenario-pack contract, asset inventory, six rubrics, FIDELITY set, pinned images, verifier schema, five starter repos with reference solutions. AGENTS.md and executive-overview.md still describe a 5-project program.          |

3. Per-project grades

| Project                      | OS | LS | Verdict                                                                                                                                                                                                                                                                                                           |
|------------------------------|----|----|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| P1 — Discovery & Diagnostics | A– | A– | Now identical across variants (LocalStack correctly deferred to P2). Strong discovery framing. But 1.2 stayed at 10 h while 1.3 was cut 18→14, so the project now runs 36 h of writing (1.1+1.4+1.5) against 24 h of hands-on. Defensible for a client-half-focused program; still the tightest technical onramp. |
| P2 — Data & Hybrid RAG       | A  | A– | Excellent scope discipline. Supplied fusion and cross-encoder wrappers de-risk 2.3, and 2.3's "justify candidate depth, fusion parameters, rerank cut-off with measured effects" converts the saved time into better learning.                                                                                    |
| P3 — Reliable AI Service     | A– | A– | Best content in the program — legacy integration, resilience, SLO lab, incident comms, and the context-switch inject. But 3.1 at 18 h now absorbs what was previously proposed as a standalone 10-hour task on top of runtime setup and provider abstraction. Highest overrun risk in the program.                |
| P4 — Security & Governance   | A  | A– | The 30-hour task is now realistic: preconfigured Keycloak realm, supplied redactor, supplied pinned CI workflow. The evidence boundary at line 274/284 is the single best-written paragraph in either document.                                                                                                   |
| P5 — Bounded Agent           | A  | A  | Evaluation design is the standout: bit-reproducible CI, judge and Ragas manual-only with three-run variance. One real risk — a ≤3B quantized model is unreliable at structured tool calling, which is exactly what 5.3 builds.                                                                                    |
| P6 — AWS Capstone & Handback | A  | A  | Byte-identical across variants. Fixed profile, prohibited resources, migration-delta record, and a required real-AWS failure demonstration. Two tight spots: 6.5 at 6 h covers three recordings; 6.6 at 8 h covers five deliverables.                                                                             |

4. What moved since the last review

| Prior blocker                           | Status                                                                    |
|-----------------------------------------|---------------------------------------------------------------------------|
| No rubrics / Must-Haves                 | Partially — declared as a build dependency; floors exist in 3 of 31 tasks |
| No hardware floor                       | Closed — lines 15–19                                                      |
| No P6 cloud profile                     | Closed — fixed profile + prohibited list                                  |
| Cost controls dropped                   | Closed — constrained role, budget actions, automated teardown             |
| GPU contradiction                       | Closed — no GPU provisioned, fine-tuning out of scope                     |
| 14-day window vs 3-week project         | Closed — endpoint by end of Week 20, 3-day SLA, one remediation           |
| P5/P6 repo continuity                   | Closed — checkpoint model                                                 |
| LocalStack ToS                          | Bounded, not closed — conditional status + no-penalty fallback            |
| LocalStack P1 fidelity violation        | Closed — LocalStack deferred to P2                                        |
| Real-AWS spend leak                     | Closed — awslocal only, client factory, CI gate                           |
| "Multi-Agent" claim                     | Closed — renamed                                                          |
| Determinism vs bots / self-set criteria | Closed — fixed transcripts, published floors                              |
| Missing legacy-integration artifact     | Closed — Task 3.1                                                         |
| No portfolio output                     | Closed — Task 6.6                                                         |
| No context-switching rep                | Closed — Task 3.4                                                         |
| Tool excess                             | Closed — ~10 services removed                                             |
| Ragas non-determinism                   | Closed — manual, variance-reported, non-gating                            |

5. Ten remaining gaps, ranked

1. Rubrics and floors don't exist. P1, P4, 6.4 and 6.5 have artifacts with no pass condition. Everything else is now blocked behind this.
2. Task 3.1 is overloaded at 18 h — runtime + provider abstraction + legacy discovery + derived spec + defensive client. Split it, or move 4 h from 3.3.
3. Admission has no production-history bar. Add "≥3 years professional production engineering with one shipped system walked through in interview."
4. ≤3B model vs LangGraph tool calling. Make the deterministic emulator the assessed agent path in P5 and Ollama optional there, or the agent lesson teaches model weakness rather than design.
5. The deployed capstone is largely TripleTen reference code (repo 5 starts from the P4 reference). Correct for grading; must be stated explicitly for the portfolio and licensing story.
6. Live rehearsal still optional (line 58/68) — the role brief names the live constraints-change drill as the exact format candidates report failing.
7. 6.5 at 6 h for three recordings and 6.6 at 8 h for five deliverables are the two tightest allocations left.
8. 6.2's required real-AWS failure demonstration needs a seeded scenario — students hold a constrained role and cannot manufacture a policy denial. Add it to the supplied-asset inventory.
9. fde-focused-6-project-tree.md is stale against both documents (old hours, old P5 name, GPU language).
10. AGENTS.md and executive-overview.md still describe a 5-project, 452-hour program with fine-tuning in P5 core.

6. Recommendation

Open-source is the primary. It now teaches slightly more than LocalStack, not less: line 231 asks students to distinguish RabbitMQ delivery semantics from the AWS adapter, and its P6 involves three genuine adapter swaps rather than an endpoint-URL change — at the same cost, with no third-party dependency.

Keep LocalStack as the documented conditional alternative. Its design is now equal, but its differentiator has shrunk: with SDK clients confined to supplied adapters, students get AWS API exposure they largely don't write. That is a thin return for a licensing condition and an extra container.

Before authoring: resolve the program-of-record question (items 9–10), then build the rubrics and floors (item 1). Everything else on the list is a 30-minute edit.

Want me to save this evaluation — and/or the earlier full review — under syllabus-fde-focused/? Your earlier claude-verdict.md request was interrupted by the model switch, so nothing was written.